

▼ Spoken Language Processing

Homework 1

1. Ссылка на занятие: <https://www.twitch.tv/deeplearningschool>
2. Материалы: <https://vk.cc/bVLCOC> <https://vk.cc/bVLD3Z>

В этом задании предлагается обучить классификатор класса возраста по голосу (пример с тем, как это можно сделать для пола см. в семинаре)

P.S. не забудьте, что если то вы работает в Colab, то вы можете поменять среду выполнения на GPU/TPU!

Вопросы по заданию/материалам: @Nestyme

```
1 !pip3 install timit-utils==0.9.0
```

```
Collecting timit-utils==0.9.0
```

```
  Downloading https://files.pythonhosted.org/packages/22/32/0c98f7f44386947b9e4080f54f09a7380c390e0b8337ab0b87050d49c43a/timit-utils-0.9.0.tar.gz
Requirement already satisfied: matplotlib in /usr/local/lib/python3.6/dist-packages (from timit-utils==0.9.0) (3.2.2)
```

```
Collecting python-speech-features
```

```
  Downloading https://files.pythonhosted.org/packages/ff/d1/94c59e20a2631985fbd2124c45177abaa9e0a4eee8ba8a305aa26fc02a8e/python-speech-features-0.6-cp36-none-any.whl
Requirement already satisfied: numpy in /usr/local/lib/python3.6/dist-packages (from timit-utils==0.9.0) (1.19.4)
```

```
Requirement already satisfied: scipy in /usr/local/lib/python3.6/dist-packages (from timit-utils==0.9.0) (1.4.1)
```

```
Requirement already satisfied: pandas in /usr/local/lib/python3.6/dist-packages (from timit-utils==0.9.0) (1.1.5)
```

```
Collecting SoundFile>=0.8.0
```

```
  Downloading https://files.pythonhosted.org/packages/eb/f2/3cbbbf3b96fb9fa91582c438b574c4ff3f45b29c772f94c400e2c99ef5db9/soundfile-0.8.0-cp36-cp37m-macosx\_10\_10\_x86\_64.whl
Requirement already satisfied: kiwisolver>=1.0.1 in /usr/local/lib/python3.6/dist-packages (from matplotlib->timit-utils==0.9.0) (1.1.0)
```

```
Requirement already satisfied: python-dateutil>=2.1 in /usr/local/lib/python3.6/dist-packages (from matplotlib->timit-utils==0.9.0) (2.6.1)
```

```
Requirement already satisfied: pyparsing!=2.0.4,!=2.1.2,!=2.1.6,>=2.0.1 in /usr/local/lib/python3.6/dist-packages (from matplotlib->timit-utils==0.9.0) (2.4.7)
```

```
Requirement already satisfied: cycler>=0.10 in /usr/local/lib/python3.6/dist-packages (from matplotlib->timit-utils==0.9.0) (0.9.0)
```

```
Requirement already satisfied: pytz>=2017.2 in /usr/local/lib/python3.6/dist-packages (from pandas->timit-utils==0.9.0) (2018.9) (2018.9)
```

```
Requirement already satisfied: cffi>=1.0 in /usr/local/lib/python3.6/dist-packages (from SoundFile>=0.8.0->timit-utils==0.9.0) (1.14.0)
```

```
Requirement already satisfied: six>=1.5 in /usr/local/lib/python3.6/dist-packages (from python-dateutil>=2.1->matplotlib->timit-utils==0.9.0) (1.14.0)
```

```
Requirement already satisfied: pyparser in /usr/local/lib/python3.6/dist-packages (from cffi>=1.0->SoundFile>=0.8.0->timit-utils==0.9.0) (1.0.0)
```

```
Building wheels for collected packages: python-speech-features
```

```
  Building wheel for python-speech-features (setup.py) ... done
```

```
  Created wheel for python-speech-features: filename=python_speech_features-0.6-cp36-none-any.whl size=5890 sha256=af11ec629012
```

Stored in directory: /root/.cache/pip/wheels/3c/42/7c/f60e9d1b40015cd69b213ad90f7c18a9264cd745b9888134be
 Successfully built python-speech-features
 Installing collected packages: python-speech-features, SoundFile, timit-utils
 Successfully installed SoundFile-0.10.3.post1 python-speech-features-0.6 timit-utils-0.9.0

1 !pip3 install torchaudio

Collecting torchaudio
 Downloading https://files.pythonhosted.org/packages/2a/f9/618434cf4e46dc975871e1516f5499abef6564ab4366f9b2321ee536be14/torchaudio-0.7.2-cp36-cp37m-macosx_10_12_x86_64.whl
 |██| 7.6MB 8.0MB/s
 Collecting torch==1.7.1
 Downloading https://files.pythonhosted.org/packages/90/4f/acf48b3a18a8f9223c6616647f0a011a5713a985336088d7c76f3a211374/torch-1.7.1-cp36-cp37m-macosx_10_12_x86_64.whl
 |██| 776.8MB 22kB/s
 Requirement already satisfied: numpy in /usr/local/lib/python3.6/dist-packages (from torch==1.7.1->torchaudio) (1.19.4)
 Requirement already satisfied: typing-extensions in /usr/local/lib/python3.6/dist-packages (from torch==1.7.1->torchaudio) (3.7.4.3)
 Requirement already satisfied: dataclasses; python_version < "3.7" in /usr/local/lib/python3.6/dist-packages (from torch==1.7.1->torchaudio) (0.8)
 ERROR: torchvision 0.8.1+cu101 has requirement torch==1.7.0, but you'll have torch 1.7.1 which is incompatible.
 Installing collected packages: torch, torchaudio
 Found existing installation: torch 1.7.0+cu101
 Uninstalling torch-1.7.0+cu101:
 Successfully uninstalled torch-1.7.0+cu101
 Successfully installed torch-1.7.1 torchaudio-0.7.2

1 ! wget <https://ndownloader.figshare.com/files/10256148>

--2020-12-27 08:45:27-- <https://ndownloader.figshare.com/files/10256148>
 Resolving ndownloader.figshare.com (ndownloader.figshare.com)... 34.248.245.115, 34.240.75.121, 3.248.141.196, ...
 Connecting to ndownloader.figshare.com (ndownloader.figshare.com)|34.248.245.115|:443... connected.
 HTTP request sent, awaiting response... 302 Found
 Location: <https://s3-eu-west-1.amazonaws.com/pfigshare-u-files/10256148/TIMIT.zip> [following]
 --2020-12-27 08:45:28-- <https://s3-eu-west-1.amazonaws.com/pfigshare-u-files/10256148/TIMIT.zip>
 Resolving s3-eu-west-1.amazonaws.com (s3-eu-west-1.amazonaws.com)... 52.218.97.10
 Connecting to s3-eu-west-1.amazonaws.com (s3-eu-west-1.amazonaws.com)|52.218.97.10|:443... connected.
 HTTP request sent, awaiting response... 200 OK
 Length: 440207227 (420M) [binary/octet-stream]
 Saving to: '10256148'

10256148 100%[=====>] 419.81M 21.7MB/s in 21s

2020-12-27 08:45:49 (20.4 MB/s) - '10256148' saved [440207227/440207227]

```
1 !unzip -q 10256148
```

```
1 import timit_utils as tu
2 import os
3 import librosa
4 import numpy as np
5 from tqdm import tqdm
6
7 import torch
8 import torch.nn as nn
9 from torch.optim import Adam
10 import torch.nn.functional as F
11
12 import matplotlib.pyplot as plt
13 from sklearn.metrics import accuracy_score
14
15 import IPython
16 _TIMIT_PATH = 'data/lisa/data/timit/raw/TIMIT'
```

▼ Задание 1

Загрузите данные для обучения. Для этого:

1. Скачайте датасет TIMIT (см семинар)
2. Соберите пары "голос" – "класс возраста" также, как на семинаре собирались пары "голос" – "пол". Аудиодорожки сконвертируйте в мелспектрограммы при помощи torchaudio либо librosa

P.S. вы можете использовать свою реализацию, а можете предложенную (см следующие ячейки)

```
1 import timit_utils as tu
```

```

2 import os
3 import librosa
4 import numpy as np
5 from tqdm import tqdm
6 import torch as t
7
8
9 class timit_dataloader:
10     def __init__(self, data_path=_TIMIT_PATH, train_mode=True, age_mode=True):
11         self.doc_file_path = os.path.join(data_path, 'DOC', 'SPKRINFO.TXT')
12         self.corpus = tu.Corpus(data_path)
13         with open(self.doc_file_path) as f:
14             self.id_age_dict = dict(
15                 [(tmp.split(' ')[0], 86 - int(tmp.split(' ')[5].split('/')[-1].replace('??', '50')))) \
16                  for tmp in f.readlines()[39:]]])
17         if train_mode:
18             self.trainset = self.create_dataset('train', age_mode=age_mode)
19             self.validset = self.create_dataset('valid', age_mode=age_mode)
20             self.testset = self.create_dataset('test', age_mode=age_mode)
21
22     def return_age(self, id):
23         return self.id_age_dict[id]
24
25     def return_data(self):
26         return self.trainset, self.validset, self.testset
27
28     def return_test(self):
29         return self.testset
30
31     def create_dataset(self, mode, age_mode=False):
32         global people
33         assert mode in ['train', 'valid', 'test']
34         if mode == 'train':
35             people = [self.corpus.train.person_by_index(i) for i in range(350)]
36         if mode == 'valid':
37             people = [self.corpus.train.person_by_index(i) for i in range(350, 400)]
38         if mode == 'test':
39             people = [self.corpus.test.person_by_index(i) for i in range(150)]

```

```

33     people = [self.corpus.test.person_by_index(i) for i in range(100)]
40     spectrograms_and_targets = []
41     for person in tqdm(people):
42         try:
43             target = self.return_age(person.name)
44             for i in range(len(person.sentences)):
45                 spectrograms_and_targets.append(
46                     self.preprocess_sample(person.sentence_by_index(i).raw_audio, target, age_mode=True))
47         except:
48             print(person.name, target)
49
50     X, y = map(np.stack, zip(*spectrograms_and_targets))
51     X = X.transpose([0, 2, 1]) # to [batch, time, channels]
52     return X, y
53
54     @staticmethod
55     def spec_to_image(spec, eps=1e-6):
56         mean = spec.mean()
57         std = spec.std()
58         spec_norm = (spec - mean) / (std + eps)
59         spec_min, spec_max = spec_norm.min(), spec_norm.max()
60         spec_scaled = 255 * (spec_norm - spec_min) / (spec_max - spec_min)
61         spec_scaled = spec_scaled.astype(np.uint8)
62         return spec_scaled
63
64     @staticmethod
65     def clasterize_by_age(age):
66         if age <= 25:
67             return 0
68         if 25 < age <= 40:
69             return 0.5
70         if age > 40:
71             return 1
72
73     def preprocess_sample(self, amplitudes, target=None, age_mode=False, sr=16000, max_length=150):
74         spectrogram = librosa.feature.melspectrogram(amplitudes, sr=sr, n_mels=128, fmin=1, fmax=8192)[: , :max_length]
75         spectrogram = np.pad(spectrogram, [[0, 0], [0, max(0, max_length - spectrogram.shape[1])]], mode='constant')
76         if target == None:

```

```

77         return np.array([self.spec_to_image(np.float32(spectrogram))]).transpose([0, 2, 1])
78     target = self.clasterize_by_age(target)
79     return self.spec_to_image(np.float32(spectrogram)), target
80
81     def preprocess_sample_inference(self, amplitudes, sr=16000, max_length=150, device='cpu'):
82         spectrogram = librosa.feature.melspectrogram(amplitudes, sr=sr, n_mels=128, fmin=1, fmax=8192)[: , :max_length]
83         spectrogram = np.pad(spectrogram, [[0, 0], [0, max(0, max_length - spectrogram.shape[1])]], mode='constant')
84         spectrogram = np.array([self.spec_to_image(np.float32(spectrogram))]).transpose([0, 2, 1])
85
86         return t.tensor(spectrogram, dtype=t.float).to(device, non_blocking=True)
87
88
89 class dataloader:
90     def __init__(self, spectrograms, targets):
91         self.data = list(zip(spectrograms, targets))
92
93     def next_batch(self, batch_size, device):
94         indices = np.random.randint(len(self.data), size=batch_size)
95
96         input = [self.data[i] for i in indices]
97
98         source = [line[0] for line in input]
99         target = [line[1] for line in input]
100
101         return self.torch_batch(source, target, device)
102
103     @staticmethod
104     def torch_batch(source, target, device):
105         # print('source',source)
106         # print('target',target)
107         return tuple(
108             [
109                 t.tensor(val, dtype=t.float).to(device, non_blocking=True)
110                 for val in [source, target]
111             ]
112         )
113
114     @staticmethod

```

```

115     def padd_sequences(lines, pad_token=0):
116         lengths = [len(line) for line in lines]
117         max_length = max(lengths)
118
119         return np.array(
120             [
121                 line + [pad_token] * (max_length - lengths[i])
122                 for i, line in enumerate(lines)
123             ]
124         )

```

Простая сверточная сеть, ее можно дотюнить или поменять по желанию

```

1
2 import torch
3 import torch.nn as nn
4 import torch.nn.functional as F
5
6
7 class Model(nn.Module):
8     def __init__(self, window_sizes=(3, 4, 5)):
9         super(Model, self).__init__()
10
11         self.convs = nn.ModuleList([
12             nn.Conv2d(1, 128, [window_size, 128], padding=(window_size - 1, 0))
13             for window_size in window_sizes
14         ])
15
16         self.fc = nn.Linear(128 * len(window_sizes), 1)
17
18     def forward(self, x):
19         # Apply a convolution + max pool layer for each window size
20         x = torch.unsqueeze(x, 1) # [B, C, T, E] Add a channel dim.
21         xs = []
22         for conv in self.convs:
23             x2 = F.relu(conv(x)) # [B, F, T, 1]

```

```

24         x2 = torch.squeeze(x2, -1) # [B, F, T]
25         x2 = F.max_pool1d(x2, x2.size(2)) # [B, F, 1]
26         xs.append(x2)
27     x = torch.cat(xs, 2) # [B, F, window]
28
29     # FC
30     x = x.view(x.size(0), -1) # [B, F * window]
31     logits = self.fc(x) # [B, class]
32     probs = torch.sigmoid(logits).view(-1)
33     return probs
34
35     def loss(self, probs, targets):
36         return nn.BCELoss()(probs.float(), targets.float())

```

```

1 device = torch.device('cuda' if torch.cuda.is_available() else 'cpu')
2 print(f'using {device} mode')
3 patience = 500
4 best_loss = 1000
5 cnt = 0
6
7 model = Model()
8 if device == torch.device('cuda'):
9     model.cuda()
10 else:
11     model.cpu()
12 model.train()

using cpu mode
Model(
  (convs): ModuleList(
    (0): Conv2d(1, 128, kernel_size=[3, 128], stride=(1, 1), padding=(2, 0))
    (1): Conv2d(1, 128, kernel_size=[4, 128], stride=(1, 1), padding=(3, 0))
    (2): Conv2d(1, 128, kernel_size=[5, 128], stride=(1, 1), padding=(4, 0))
  )
  (fc): Linear(in_features=384, out_features=1, bias=True)
)

```



```

1 _timit_dataloader = timit_dataloader()
2 train, valid, test = _timit_dataloader.return_data()
3
4 trainset = dataloader(*train)
5 validset = dataloader(*valid)
6 testset = dataloader(*test)
7 BATCH_SIZE = 64
8
9 optimizer = Adam(
10     [p for p in model.parameters() if p.requires_grad], betas=(0.9, 0.999), eps=1e-5
11 )

```

```

100%|██████████| 350/350 [00:42<00:00, 8.28it/s]
100%|██████████| 50/50 [00:06<00:00, 8.28it/s]
100%|██████████| 150/150 [00:18<00:00, 8.03it/s]

```

▼ Задание 2

1. Обучите свой классификатор категории возраста
2. Попробуйте улучшить результат. Можно попробовать усложнить сетку, подвигать границы категорий, поискать новые данные, что угодно, кроме учиться на тесте :)

Чем больше границы тем лучше качество , чем меньше границы тем хуже качество.

3. Какой подход оказался самым эффективным? Как думаете, почему?

Эффективней всего обучать маленьких границах но при этом предсказывать более большие тогда качество должно повыситься
Возможно вообще надо предсказывать в днях или месяцах ,а потом переводить в года. Не успел проверить.

4. Как считаете, где можно было бы применить такой классификатор в качестве вспомогательной задачи?

Как вспомогательная задача может быть применима в голосовых помощниках, или умных домах

```

1 import torch as t
2 for i in tqdm(range(100)):

```

```

3
4 optimizer.zero_grad()
5 # print('batch',BATCH_SIZE,'device',device)
6 input, target = trainset.next_batch(BATCH_SIZE, device=device)
7 out = model(input)
8 loss = model.loss(out, target)
9 loss.backward()
10 optimizer.step()
11
12 if i % 50 == 0:
13     model.eval()
14
15     with torch.no_grad():
16         optimizer.zero_grad()
17
18         input, target = validset.next_batch(BATCH_SIZE, device=device)
19         out = model(input)
20         valid_loss = model.loss(out, target)
21         out, target = out.cpu().detach().numpy(), target.cpu().detach().numpy()
22         # print(out, target)
23         new_out = np.array([])
24         for tmp in out:
25             if tmp <= 0.33:
26                 new_out = np.append(new_out, 1)
27             # elif tmp <= 0.66:
28             #     new_out = np.append(new_out, 2)
29             else:
30                 new_out = np.append(new_out, 3)
31         # out = [1. if tmp > 0.5 else 0 for tmp in out]
32         target = np.where(target==1, 2, target)
33         target = np.where(target==0.5, 1, target)
34         target = np.where(target==0, 0, target)
35         # print(new_out)
36         # print(target)
37         print(f'accuracy_score:{accuracy_score(new_out, target)}')
38         # print(f'accuracy_score:{clf.score(new_out, target)}')
39         print("i {}, valid {}".format(i, valid_loss.item()))
40         # print(" ")

```

```

40         print( '_____ ' )
41
42     model.train()
43
44     if i % 50 == 0 and best_loss > valid_loss.item():
45         best_loss = valid_loss.item()
46         cnt = 0
47     else:
48         cnt += 1
49
50     if cnt > patience:
51         break
52 print('training finished')
53

```

```

31%|██████| 31/100 [00:14<00:30, 2.25it/s]batch 64 device cpu
32%|██████| 32/100 [00:14<00:30, 2.26it/s]batch 64 device cpu
33%|██████| 33/100 [00:14<00:29, 2.27it/s]batch 64 device cpu
34%|██████| 34/100 [00:15<00:29, 2.24it/s]batch 64 device cpu
35%|██████| 35/100 [00:15<00:29, 2.23it/s]batch 64 device cpu
36%|██████| 36/100 [00:16<00:28, 2.24it/s]batch 64 device cpu
37%|██████| 37/100 [00:16<00:28, 2.22it/s]batch 64 device cpu
38%|██████| 38/100 [00:17<00:27, 2.24it/s]batch 64 device cpu
39%|██████| 39/100 [00:17<00:27, 2.24it/s]batch 64 device cpu
40%|██████| 40/100 [00:18<00:26, 2.25it/s]batch 64 device cpu
41%|██████| 41/100 [00:18<00:26, 2.25it/s]batch 64 device cpu
42%|██████| 42/100 [00:18<00:25, 2.26it/s]batch 64 device cpu
43%|██████| 43/100 [00:19<00:25, 2.27it/s]batch 64 device cpu
44%|██████| 44/100 [00:19<00:24, 2.27it/s]batch 64 device cpu
45%|██████| 45/100 [00:20<00:24, 2.27it/s]batch 64 device cpu
46%|██████| 46/100 [00:20<00:23, 2.25it/s]batch 64 device cpu
47%|██████| 47/100 [00:21<00:23, 2.26it/s]batch 64 device cpu
48%|██████| 48/100 [00:21<00:23, 2.26it/s]batch 64 device cpu
49%|██████| 49/100 [00:22<00:22, 2.22it/s]batch 64 device cpu
50%|██████| 50/100 [00:22<00:22, 2.24it/s]batch 64 device cpu
51%|██████| 51/100 [00:23<00:26, 1.88it/s]accuracy_score:0.296875
i 50, valid 0.9732868671417236

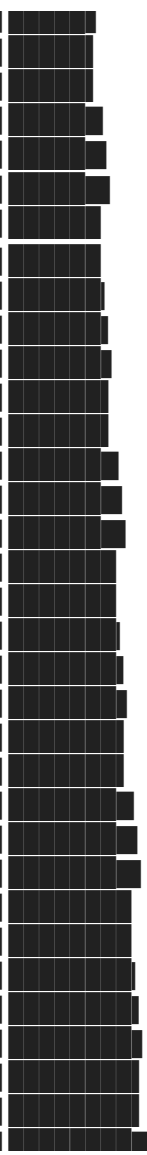
```

batch 64 device cpu

```

52%|██████| 52/100 [00:23<00:24, 1.98it/s]batch 64 device cpu
53%|██████| 53/100 [00:24<00:22, 2.07it/s]batch 64 device cpu

```



```

54%|██████████| 54/100 [00:24<00:21, 2.12it/s]batch 64 device cpu
55%|██████████| 55/100 [00:24<00:20, 2.16it/s]batch 64 device cpu
56%|██████████| 56/100 [00:25<00:20, 2.18it/s]batch 64 device cpu
57%|██████████| 57/100 [00:25<00:20, 2.14it/s]batch 64 device cpu
58%|██████████| 58/100 [00:26<00:19, 2.16it/s]batch 64 device cpu
59%|██████████| 59/100 [00:26<00:18, 2.17it/s]batch 64 device cpu
60%|██████████| 60/100 [00:27<00:18, 2.19it/s]batch 64 device cpu
61%|██████████| 61/100 [00:27<00:17, 2.19it/s]batch 64 device cpu
62%|██████████| 62/100 [00:28<00:17, 2.20it/s]batch 64 device cpu
63%|██████████| 63/100 [00:28<00:16, 2.20it/s]batch 64 device cpu
64%|██████████| 64/100 [00:29<00:16, 2.22it/s]batch 64 device cpu
65%|██████████| 65/100 [00:29<00:15, 2.20it/s]batch 64 device cpu
66%|██████████| 66/100 [00:29<00:15, 2.21it/s]batch 64 device cpu
67%|██████████| 67/100 [00:30<00:14, 2.21it/s]batch 64 device cpu
68%|██████████| 68/100 [00:30<00:14, 2.19it/s]batch 64 device cpu
69%|██████████| 69/100 [00:31<00:14, 2.21it/s]batch 64 device cpu
70%|██████████| 70/100 [00:31<00:13, 2.22it/s]batch 64 device cpu
71%|██████████| 71/100 [00:32<00:12, 2.25it/s]batch 64 device cpu
72%|██████████| 72/100 [00:32<00:12, 2.24it/s]batch 64 device cpu
73%|██████████| 73/100 [00:33<00:12, 2.24it/s]batch 64 device cpu
74%|██████████| 74/100 [00:33<00:11, 2.24it/s]batch 64 device cpu
75%|██████████| 75/100 [00:33<00:11, 2.25it/s]batch 64 device cpu
76%|██████████| 76/100 [00:34<00:10, 2.26it/s]batch 64 device cpu
77%|██████████| 77/100 [00:34<00:10, 2.24it/s]batch 64 device cpu
78%|██████████| 78/100 [00:35<00:09, 2.26it/s]batch 64 device cpu
79%|██████████| 79/100 [00:35<00:09, 2.26it/s]batch 64 device cpu
80%|██████████| 80/100 [00:36<00:08, 2.26it/s]batch 64 device cpu
81%|██████████| 81/100 [00:36<00:08, 2.26it/s]batch 64 device cpu
82%|██████████| 82/100 [00:37<00:07, 2.28it/s]batch 64 device cpu
83%|██████████| 83/100 [00:37<00:07, 2.28it/s]batch 64 device cpu
84%|██████████| 84/100 [00:37<00:07, 2.27it/s]batch 64 device cpu
85%|██████████| 85/100 [00:38<00:06, 2.27it/s]batch 64 device cpu
86%|██████████| 86/100 [00:38<00:06, 2.25it/s]batch 64 device cpu
87%|██████████| 87/100 [00:39<00:05, 2.26it/s]batch 64 device cpu

```

```

1 model.eval()
2
3 def predict(wavfile):
4     waveform, _ = librosa.load(wavfile, sr=16000)
5
6     input = _timit_dataloader.preprocess_sample(waveform)
7     with torch.no_grad():

```

```

7     with torch.no_grad():
8         out = model(torch.tensor(input, dtype=torch.float).to(device))
9         out = out.cpu().detach().numpy()
10    print(out)
11    if out <= 0.33:
12        out = '<25'
13    elif out <= 0.66:
14        out = '>25<40'
15    else:
16        out = '>40'
17    # out = 'female' if out < 0.5 else 'male'
18    return out

```

```

1 # Code for recording audio from the browser
2 from IPython.display import Javascript
3 from google.colab import output
4 from base64 import b64decode
5 import IPython
6 import uuid
7 from google.colab import output
8
9
10 class InvokeButton(object):
11     def __init__(self, title, callback):
12         self._title = title
13         self._callback = callback
14
15     def _repr_html_(self):
16         from google.colab import output
17         callback_id = 'button-' + str(uuid.uuid4())
18         output.register_callback(callback_id, self._callback)
19
20         template = """<button id="{callback_id}" style="cursor:pointer;background-color:#EEEEEE;border-color:#E0E0E0;padding:5px 15px.
21             <script>
22                 document.querySelector("#{callback_id}").onclick = (e) => {{
23                     google.colab.kernel.invokeFunction('{callback_id}', [], {{}})}
24                     e.preventDefault();
25             ""

```

```

26     </script>""
27     html = template.format(title=self._title, callback_id=callback_id)
28     return html
29
30 RECORD = ""
31 const sleep = time => new Promise(resolve => setTimeout(resolve, time))
32 const b2text = blob => new Promise(resolve => {
33     const reader = new FileReader()
34     reader.onloadend = e => resolve(e.srcElement.result)
35     reader.readAsDataURL(blob)
36 })
37 var record = time => new Promise(async resolve => {
38     stream = await navigator.mediaDevices.getUserMedia({ audio: true })
39     recorder = new MediaRecorder(stream)
40     chunks = []
41     recorder.ondataavailable = e => chunks.push(e.data)
42     recorder.start()
43     await sleep(time)
44     recorder.onstop = async ()=>{
45         blob = new Blob(chunks)
46         text = await b2text(blob)
47         resolve(text)
48     }
49     recorder.stop()
50 })
51 ""
52
53 def record(sec=3):
54     display(Javascript(RECORD))
55     s = output.eval_js('record(%d)' % (sec*1000))
56     b = b64decode(s.split(',')[1])
57     with open('audio.wav', 'wb+') as f:
58         f.write(b)
59     return 'audio.wav'

```

```

1 def classify():
2     print("Now recording for 3 seconds say what you will")

```

```
2 print('Now recording for 3 seconds, say what you will...')
3 record()
4 os.system('ffmpeg -i audio.wav -ar 16000 -y audio.wav')
5 print(f"Audio recording complete, guess it is {predict('audio.wav')}")
6
7 InvokeButton('Start recording', classify)
```

Start recording

```
Now recording for 3 seconds, say what you will...
[0.00109952]
Audio recording complete, guess it is <25
Now recording for 3 seconds, say what you will...
[7.352042e-05]
Audio recording complete, guess it is <25
Now recording for 3 seconds, say what you will...
[0.00268975]
Audio recording complete, guess it is <25
Now recording for 3 seconds, say what you will...
[0.03905485]
Audio recording complete, guess it is <25
Now recording for 3 seconds, say what you will...
[0.00256719]
Audio recording complete, guess it is <25
```

```
1 IPython.display.Audio('audio.wav')
```

0:03 / 0:03

1

